

PHYSICS-INFORMED TRAINING AT THE OPERATOR LEVEL

Anonymous authors

Paper under double-blind review

ABSTRACT

We present an approach to train neural operators physics-informed, i.e., in a self-supervised fashion, purely relying on the governing physical equations. This enables training on infinite datasets where no sample is visited twice and removes the need to produce training data in an offline phase. While previous attempts for physics-informed training at the operator level were limited to benign problems and performed considerably worse than data-driven learning, our approach matches the gold-standard of data-driven training at a similar optimization cost. Our method is based on modifying the physics-informed loss functions to include a preconditioner such as a multigrid method to drastically reduce the ill-conditioning caused by unbounded PDE operators. We derive the improved conditioning theoretically, and illustrate the potential of the method for the Poisson, Allen-Cahn, and stationary Stokes equation, where it performs virtually identical to data-driven training at a comparable computational cost.

1 INTRODUCTION

Neural Operators Despite a century of development of numerical methods, partial differential equations (PDEs) remain expensive to solve, even more so in many-query situations, such as uncertainty quantification, optimal control, and inverse problems. The development of fast PDE solution techniques, e.g., through reduced-order/surrogate models therefore is of utmost importance. In the wake of the deep learning revolution, neural operators emerged as a powerful computational tool for surrogate modeling, drastically reducing PDE solution time in many applications. Neural operators approximate maps between infinite-dimensional spaces (operators) by a neural network, and are typically trained in a supervised fashion. They hence require a dataset that is built in an offline fashion, before they can be trained, and eventually used downstream. Data generation can constitute a substantial part of the computation and engineering effort, even more so in the context of physics foundation models that are pre-trained over an array of different PDE instances.

Physics-Informed Training Whenever the governing physics in the form of PDEs are known, neural operators can—in principle—be trained in a self-supervised fashion, relying on the physics residual or variational principles as a loss function. This approach is referred to as *physics-informed neural operator training*, and a natural idea present in the community since the inceptions of neural operators []. However, the success of physics-informed training of neural operators has remained limited to benign equations and learning tasks. More often than not, purely physics-informed learning fails to converge. In fact, supervised learning with its offline data-generation pipeline remains the gold standard for almost all applications. Nevertheless, physics-informed training offers unique




Figure 1: Adam on the plain versus the preconditioned Poisson least-squares loss, $N \times N$ grid. Placeholder numbers.

advantages, such as the training in the infinite data-limit, where no sample is revisited, and, of course, it removes the need of an expensive dataset generation process.

Ill-Conditioning of PDE Operators In this manuscript, we show that the pathologies of physics-informed neural operator learning can be attributed to a single foundational property of PDE operators: PDE operators are unbounded maps, meaning they magnify their inputs without bounds, and thus lead to severe ill-conditioning in physics-informed loss functions. To make this concrete, take the derivative operator

$$\frac{d}{dx} : C^1(0, 1) \subset L^2(0, 1) \rightarrow L^2(0, 1), \quad u \mapsto u', \quad (1)$$

which amplifies oscillations: $x \mapsto \sin(\omega x)$ is bounded in $L^2(0, 1)$ independent of $\omega \in \mathbb{R}$, but its derivative $x \mapsto \omega \cos(\omega x)$ grows without bound in $L^2(0, 1)$ when $\omega \rightarrow \infty$. This phenomenon is inherent in all PDE operators. In the case of Poisson’s equation, and for a linear model that is the linear combination of a given basis set $\phi_1, \dots, \phi_n$, the physics-informed least-squares loss¹ on a single data point f takes the form

$$L_{\text{LS}}(u) = \frac{1}{2} \|Au - f\|^2, \quad A_{ij} = \int_{\Omega} \nabla \phi_j \nabla \phi_i \, dx, \quad (2)$$

where $A \in \mathbb{R}^{N \times N}$ is the stiffness matrix. The condition number of A scales like h^{-2} for standard bases, where h is the characteristic mesh-size. Conditioning thus deteriorates under mesh refinement, as higher oscillating frequencies can be resolved, the same manifestation of the unbounded nature of PDE operators stressed in 1. Note that $\nabla^2 L_{\text{LS}} = A^2$, hence the condition number governing the convergence speed of gradient descent (GD) scales like h^{-4} . Figure 1 illustrates that this ill-conditioning is so severe that training with first-order methods like GD and Adam becomes infeasible already for very coarse meshes, even in the benign case of a linear ansatz (no neural network at all). In the remainder of this article, we show that precisely this ill-conditioning renders physics-informed neural operator learning infeasible.

Preconditioned Physics-Informed Training To counter the ill-conditioning of the least-squares loss equation 2, we rely on *preconditioning*—precisely the tool the numerical analysis community uses to efficiently employ iterative methods in ill-conditioned settings in classical scientific computing. Abstractly, we *modify the least-squares loss* to include an invertible preconditioner P and a positive definite and symmetric weighting matrix B

$$L_{\text{PLS}}(u) = \frac{1}{2} \|P[Au - f]\|_B^2, \quad (3)$$

We stress that this construction is agnostic to the neural operator architecture employed and adds zero overhead to the inference once training finishes as it only influences the loss function. On a conceptual level, the preconditioned loss equation 3 changes the norm in which the error is measured. For any given P we can recast the

$$L_{\text{PLS}}(u) = \frac{1}{2} \|P[Au - f]\|_B^2 = \frac{1}{2} (PAe, PAe)_B$$

with the error $e = u - A^{-1}f$. This means P maps Ae into a different space and thereby improves conditioning, and B measures the error in this space in an appropriate manner. The choice $P = A^{-1}$ and $B = I^2$ recovers data-driven training and reduces the condition number of $\nabla^2 L_{\text{PLS}}$ to unity. In practice, this is realized through fast approximate inverses $P \approx A^{-1}$, e.g., via multigrid methods []. Note that $P \approx A^{-1}$ is not the only choice; for saddle point problems, we alternatively employ block-diagonal preconditioners. Additionally, whenever applicable, we investigate energy formulations, as these are naturally less ill-conditioned.

1.1 MAIN CONTRIBUTIONS

The main contributions are

¹If the equation admits, energy formulations (variational principles) are an attractive alternative.

²Or alternatively, $B = M$, where M is the mass matrix.

- **Efficient physics-informed neural operator training.** We demonstrate that preconditioned physics-informed neural operator training allows data-free training of neural operators at comparable cost and accuracy of the gold standard of data-driven training. Our approach vastly outperforms previous approaches and is demonstrated for the Poisson, Allen-Cahn, and Stokes equation.
- **Flexibility and inference speed.** Our method can be combined with any neural operator architecture and once trained, adds no additional cost to the inference.
- **Training at infinite data limit.** Physics-informed training allows to practically approximate the infinite data limit, where no sample is visited twice. We demonstrate empirically that this is a favorable training paradigm for neural operators.

1.2 RELATED LITERATURE

- Discussion of physics-informed training in general, in particular for PINNs and NQS.
- Paper of Wolfgang and Markus on preconditioning.
- OG paper on physics-informed neural operators.
- Papers on suboptimal fixes for physics-informed neural operator training.

2 PHYSICS-INFORMED TRAINING OF NEURAL OPERATORS

Our goal is to approximate an operator $\mathcal{G} : \mathcal{F} \rightarrow \mathcal{U}$ between function spaces $\mathcal{U}$ and $\mathcal{F}$. We aim to approximate $\mathcal{G}$ by a neural network $\mathcal{G} \approx \mathcal{G}_\theta$. Neural operators are a special class of neural networks well-suited to the infinite-dimensional nature of the function spaces appearing in PDE applications. However, the concrete neural network structure of $\mathcal{G}_\theta$ is secondary for our proposed training methodology. A simple but crucial observation we make, is that we can always write the neural operator output in the *interpolated form*

$$\mathcal{G}_\theta(f) = \sum_{i=1}^n \mathbf{u}(\theta, f)_i \phi_i, \quad (4)$$

where $\mathbf{u}(\theta, f) \in \mathbb{R}^n$ are coefficients depending on the trainable weights and the functional input f . The functions $\phi_1, \dots, \phi_n$ are a fixed basis and span a finite-dimensional function space $\mathcal{U}_h$, where h denotes a characteristic mesh-size³. Typical choices are finite element spaces, like piecewise polynomial spaces on a triangulated computational domain. It is crucial to stress that $\phi_1, \dots, \phi_n$, and thus $\mathcal{U}_h$ is a modeling choice, and a standard neural operator architecture will not automatically map into $\mathcal{U}_h$. However, we can always interpolate into $\mathcal{U}_h$. This is a typically sparse, non-trainable operation that adds negligible computational overhead and we will from now on assume that our neural operator is the form described in equation 4.

2.1 PHYSICS-INFORMED LOSS FUNCTIONS

For simplicity of exposition, we will work in the setting of linear PDEs that can be formulated in terms of bilinear forms, and we will restrict the discussion to learning the solution operator of such PDEs, although our experimental evaluation is broader. The operator learning problem is to find $u \in \mathcal{U}$ such that for given $f \in \mathcal{F}$ it holds

$$a(u, v) = f(v), \quad \text{for all } v \in \mathcal{U}. \quad (5)$$

Here, a is a bounded and symmetric bilinear form on $\mathcal{U}$, and the setting implies $\mathcal{F} = \mathcal{U}^*$. We assume all spaces to be Hilbert spaces. The solution operator is $\mathcal{G} : f \mapsto u$, where u solves equation 5. The weak equation can be written as a matrix equation, once u is approximated in $\mathcal{U}_h$, and takes the form

$$\mathbf{A}\mathbf{u} = \mathbf{f}, \quad \text{with } A_{ij} = a(\phi_j, \phi_i), \quad \mathbf{f}_i = f(\phi_i).$$

From this equation, we derive the single sample physics-informed loss function

$$L_{\text{LS}}(\theta) = \frac{1}{2} \|\mathbf{A}\mathbf{u}(\theta, f) - \mathbf{f}\|^2. \quad (6)$$

³Or more general, a quantification of how “fine” the space $\mathcal{U}_h$ is.

For analysis purposes, we will frequently ignore that nonlinear parametrization $\mathbf{u}(\theta, f)$ and work in the Euclidean space of $\mathbf{u} \in \mathbb{R}^n$ directly. We indicate this by writing $L_{LS}(u) = \frac{1}{2} \|A\mathbf{u} - \mathbf{f}\|^2$ with the understanding that everything is in the variable $\mathbf{u}$. If a is coercive and symmetric an energy formulation is an attractive alternative to equation 5. In this case, we can find the solution u via a variational principle as the minimizer of the energy $\frac{1}{2}a(v, v) - f(v)$, or in terms of the neural operator

$$L_{\text{Var}}(\theta) = \frac{1}{2} \mathbf{u}(\theta, f)^\top A \mathbf{u}(\theta, f) - \mathbf{f}^\top \mathbf{u}(\theta, f) \quad (7)$$

and we can directly utilize L_{Var} to train the neural operator.

2.2 PRECONDITIONING

3 NUMERICAL RESULTS

3.1 POISSON EQUATION

3.2 ALLEN-CAHN EQUATION

3.3 STOKES EQUATION

3.4 STYLE

Papers to be submitted to ICLR 2027 must be prepared according to the instructions presented here.

Authors are required to use the ICLR $\text{\LaTeX}$ style files obtainable at the ICLR website. Please make sure you use the current files and not previous versions. Tweaking the style files may be grounds for rejection.

3.5 RETRIEVAL OF STYLE FILES

The style files for ICLR and other conference information are available online at:

<http://www.iclr.cc/>

The file `iclr2027_conference.pdf` contains these instructions and illustrates the various formatting requirements your ICLR paper must satisfy. Submissions must be made using $\text{\LaTeX}$ and the style files `iclr2027_conference.sty` and `iclr2027_conference.bst` (to be used with $\text{\LaTeX}2\epsilon$). The file `iclr2027_conference.tex` may be used as a “shell” for writing your paper. All you have to do is replace the author, title, abstract, and text of the paper with your own.

The formatting instructions contained in these style files are summarized in sections 4, 5, and 6 below.

4 GENERAL FORMATTING INSTRUCTIONS

The text must be confined within a rectangle 5.5 inches (33 picas) wide and 9 inches (54 picas) long. The left margin is 1.5 inch (9 picas). Use 10 point type with a vertical spacing of 11 points. Times New Roman is the preferred typeface throughout. Paragraphs are separated by 1/2 line space, with no indentation.

Paper title is 17 point, in small caps and left-aligned. All pages should start at 1 inch (6 picas) from the top of the page.

Authors’ names are set in boldface, and each name is placed above its corresponding address. The lead author’s name is to be listed first, and the co-authors’ names are set to follow. Authors sharing the same address can be on the same line.

Please pay special attention to the instructions in section 6 regarding figures, tables, acknowledgments, and references.

There will be a strict upper limit of **9 pages** for the main text of the initial submission, with unlimited additional pages for citations. This limit will be expanded to **10 pages** for rebuttal/camera ready.

5 HEADINGS: FIRST LEVEL

First level headings are in small caps, flush left and in point size 12. One line space before the first level heading and 1/2 line space after the first level heading.

5.1 HEADINGS: SECOND LEVEL

Second level headings are in small caps, flush left and in point size 10. One line space before the second level heading and 1/2 line space after the second level heading.

5.1.1 HEADINGS: THIRD LEVEL

Third level headings are in small caps, flush left and in point size 10. One line space before the third level heading and 1/2 line space after the third level heading.

6 CITATIONS, FIGURES, TABLES, REFERENCES

These instructions apply to everyone, regardless of the formatter being used.

6.1 CITATIONS WITHIN THE TEXT

Citations within the text should be based on the `natbib` package and include the authors' last names and year (with the "et al." construct for more than two authors). When the authors or the publication are included in the sentence, the citation should not be in parenthesis using `\citet{}` (as in "See Hinton et al. (2006) for more information."). Otherwise, the citation should be in parenthesis using `\citep{}` (as in "Deep learning shows promise to make progress towards AI (Bengio & LeCun, 2007).").

The corresponding references are to be listed in alphabetical order of authors, in the REFERENCES section. As to the format of the references themselves, any style is acceptable as long as it is used consistently.

6.2 FOOTNOTES

Indicate footnotes with a number⁴ in the text. Place the footnotes at the bottom of the page on which they appear. Precede the footnote with a horizontal rule of 2 inches (12 picas).⁵

6.3 FIGURES

All artwork must be neat, clean, and legible. Lines should be dark enough for purposes of reproduction; art work should not be hand-drawn. The figure number and caption always appear after the figure. Place one line space before the figure caption, and one line space after the figure. The figure caption is lower case (except for first word and proper nouns); figures are numbered consecutively.

Make sure the figure caption does not get separated from the figure. Leave sufficient space to avoid splitting the figure and figure caption.

You may use color figures. However, it is best for the figure captions and the paper body to make sense if the paper is printed either in black/white or in color.

⁴Sample of the first footnote

⁵Sample of the second footnote




Figure 2: Sample figure caption.

Table 1: Sample table title

PART	DESCRIPTION
Dendrite	Input terminal
Axon	Output terminal
Soma	Cell body (contains cell nucleus)

6.4 TABLES

All tables must be centered, neat, clean and legible. Do not use hand-drawn tables. The table number and title always appear before the table. See Table 1.

Place one line space before the table title, one line space after the table title, and one line space after the table. The table title must be lower case (except for first word and proper nouns); tables are numbered consecutively.

7 DEFAULT NOTATION

In an attempt to encourage standardized notation, we have included the notation file from the textbook, *Deep Learning* Goodfellow et al. (2016) available at https://github.com/goodfeli/dlbook_notation/. Use of this style is not required and can be disabled by commenting out `math_commands.tex`.

Numbers and Arrays

a	A scalar (integer or real)
$\mathbf{a}$	A vector
$\mathbf{A}$	A matrix
$\mathbf{A}$	A tensor
$\mathbf{I}_n$	Identity matrix with n rows and n columns
$\mathbf{I}$	Identity matrix with dimensionality implied by context
$\mathbf{e}^{(i)}$	Standard basis vector $[0, \dots, 0, 1, 0, \dots, 0]$ with a 1 at position i
$\text{diag}(\mathbf{a})$	A square, diagonal matrix with diagonal entries given by $\mathbf{a}$
$\mathbf{a}$	A scalar random variable
$\mathbf{a}$	A vector-valued random variable
$\mathbf{A}$	A matrix-valued random variable

Sets and Graphs

$\mathbb{A}$	A set
$\mathbb{R}$	The set of real numbers
$\{0, 1\}$	The set containing 0 and 1
$\{0, 1, \dots, n\}$	The set of all integers between 0 and n
$[a, b]$	The real interval including a and b
$(a, b]$	The real interval excluding a but including b
$\mathbb{A} \setminus \mathbb{B}$	Set subtraction, i.e., the set containing the elements of $\mathbb{A}$ that are not in $\mathbb{B}$
$\mathcal{G}$	A graph
$Pa_{\mathcal{G}}(x_i)$	The parents of x_i in $\mathcal{G}$

Indexing

a_i	Element i of vector $\mathbf{a}$, with indexing starting at 1
$\mathbf{a}_{-i}$	All elements of vector $\mathbf{a}$ except for element i
$A_{i,j}$	Element i, j of matrix $\mathbf{A}$
$\mathbf{A}_{i,:}$	Row i of matrix $\mathbf{A}$
$\mathbf{A}_{:,i}$	Column i of matrix $\mathbf{A}$
$\mathbf{A}_{i,j,k}$	Element (i, j, k) of a 3-D tensor $\mathbf{A}$
$\mathbf{A}_{::,i}$	2-D slice of a 3-D tensor
$\mathbf{a}_i$	Element i of the random vector $\mathbf{a}$

Calculus

$\frac{dy}{dx}$	Derivative of y with respect to x
$\frac{\partial y}{\partial x}$	Partial derivative of y with respect to x
$\nabla_{\mathbf{x}} y$	Gradient of y with respect to $\mathbf{x}$
$\nabla_{\mathbf{X}} y$	Matrix derivatives of y with respect to $\mathbf{X}$
$\nabla_{\mathbf{X}} y$	Tensor containing derivatives of y with respect to $\mathbf{X}$
$\frac{\partial f}{\partial \mathbf{x}}$	Jacobian matrix $\mathbf{J} \in \mathbb{R}^{m \times n}$ of $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$
$\nabla_{\mathbf{x}}^2 f(\mathbf{x})$ or $\mathbf{H}(f)(\mathbf{x})$	The Hessian matrix of f at input point $\mathbf{x}$
$\int f(\mathbf{x}) d\mathbf{x}$	Definite integral over the entire domain of $\mathbf{x}$
$\int_{\mathbb{S}} f(\mathbf{x}) d\mathbf{x}$	Definite integral with respect to $\mathbf{x}$ over the set $\mathbb{S}$

Probability and Information Theory

378	$P(a)$	A probability distribution over a discrete variable
379	$p(a)$	A probability distribution over a continuous variable, or
380		over a variable whose type has not been specified
381		
382	$a \sim P$	Random variable a has distribution P
383	$\mathbb{E}_{x \sim P}[f(x)]$ or $\mathbb{E}f(x)$	Expectation of $f(x)$ with respect to $P(x)$
384	$\text{Var}(f(x))$	Variance of $f(x)$ under $P(x)$
385	$\text{Cov}(f(x), g(x))$	Covariance of $f(x)$ and $g(x)$ under $P(x)$
386	$H(x)$	Shannon entropy of the random variable x
387	$D_{\text{KL}}(P Q)$	Kullback-Leibler divergence of P and Q
388	$\mathcal{N}(x; \mu, \Sigma)$	Gaussian distribution over x with mean μ and covariance
389		Σ
390		
391		
392		

Functions

394	$f : \mathbb{A} \rightarrow \mathbb{B}$	The function f with domain $\mathbb{A}$ and range $\mathbb{B}$
395	$f \circ g$	Composition of the functions f and g
396	$f(x; \theta)$	A function of x parametrized by θ . (Sometimes we write
397		$f(x)$ and omit the argument θ to lighten notation)
398		
399	$\log x$	Natural logarithm of x
400	$\sigma(x)$	Logistic sigmoid, $\frac{1}{1 + \exp(-x)}$
401	$\zeta(x)$	Softplus, $\log(1 + \exp(x))$
402	$\ x\ _p$	L^p norm of x
403	$\ x\ $	L^2 norm of x
404	x^+	Positive part of x , i.e., $\max(0, x)$
405	$\mathbf{1}_{\text{condition}}$	is 1 if the condition is true, 0 otherwise
406		
407		
408		
409		

8 FINAL INSTRUCTIONS

Do not change any aspects of the formatting parameters in the style files. In particular, do not modify the width or length of the rectangle the text should fit into, and do not change font sizes (except perhaps in the REFERENCES section; see below). Please note that pages should be numbered.

9 PREPARING POSTSCRIPT OR PDF FILES

Please prepare PostScript or PDF files with paper size “US Letter”, and not, for example, “A4”. The `-t letter` option on `dvips` will produce US Letter files.

Consider directly generating PDF files using `pdflatex` (especially if you are a MiKTeX user). PDF figures must be substituted for EPS figures, however.

Otherwise, please generate your PostScript and PDF files with the following commands:

```
dvips mypaper.dvi -t letter -Ppdf -G0 -o mypaper.ps
ps2pdf mypaper.ps mypaper.pdf
```

9.1 MARGINS IN LATEX

Most of the margin problems come from figures positioned by hand using `\special` or other commands. We suggest using the command `\includegraphics` from the `graphicx` package.

Always specify the figure width as a multiple of the line width as in the example below using .eps graphics

```
\usepackage[dvips]{graphicx} ...
\includegraphics[width=0.8\linewidth]{myfile.eps}
```

or

```
\usepackage[pdftex]{graphicx} ...
\includegraphics[width=0.8\linewidth]{myfile.pdf}
```

for .pdf graphics. See section 4.4 in the graphics bundle documentation (<http://www.ctan.org/tex-archive/macros/latex/required/graphics/grfguide.ps>)

A number of width problems arise when LaTeX cannot properly hyphenate a line. Please give LaTeX hyphenation hints using the \- command.

AI USE STATEMENT

(This section is **required** and does not count toward the page limit.)

In this work, we used generative AI tools for [tasks with required disclosure]. We have not used generative AI tools for [other tasks with required disclosure], and [the rest of the required disclosure tasks] are not applicable to this work. Additionally, we used generative AI tools for [tasks with recommended disclosure]. We have reviewed all AI-assisted work. [Elaborate. For example, “we checked LLM-generated research ideas for potential plagiarism through a manual literature survey”, “LLM-generated code was verified and tested for correctness by 2 authors”, etc.]. We take responsibility for the final content of this work, including text, claims or artifacts produced with the aid of generative AI.

See the ICLR 2027 AI Policy for Authors for more details. This statement should not be more than 1 page.

ETHICS STATEMENT

(This section is **recommended** and does not count toward the page limit.)

If authors feel that their paper submission raises questions regarding the Code of Ethics, they are encouraged to include a paragraph of Ethics Statement (at the end of the main text before references) to address potential concerns where appropriate. Topics include, but are not limited to, studies that involve human subjects, practices to data set releases, potentially harmful insights, methodologies and applications, potential conflicts of interest and sponsorship, discrimination/bias/fairness concerns, privacy and security issues, legal compliance, and research integrity issues (e.g., IRB, documentation, research ethics). This statement should not be more than 1 page.

REPRODUCIBILITY STATEMENT

(This section is **recommended** and does not count toward the page limit.)

It is important that the work published in ICLR is reproducible. Authors are strongly encouraged to include a paragraph-long Reproducibility Statement at the end of the main text (before references) to discuss the efforts that have been made to ensure reproducibility. This paragraph should not itself describe details needed for reproducing the results, but rather reference the parts of the main paper, appendix, and supplemental materials that will help with reproducibility. For example, for novel models or algorithms, a link to an anonymous downloadable source code can be submitted as supplementary materials; for theoretical results, clear explanations of any assumptions and a complete proof of the claims can be included in the appendix; for any datasets used in the experiments, a complete description of the data processing steps can be provided in the supplementary materials. Each of the above are examples of things that can be referenced in the reproducibility statement.

AUTHOR CONTRIBUTIONS

If you'd like to, you may include a section for author contributions as is done in many journals. This is optional and at the discretion of the authors.

ACKNOWLEDGMENTS

Use unnumbered third level headings for the acknowledgments. All acknowledgments, including those to funding agencies, go at the end of the paper.

REFERENCES

- Yoshua Bengio and Yann LeCun. Scaling learning algorithms towards AI. In *Large Scale Kernel Machines*. MIT Press, 2007.
- Ian Goodfellow, Yoshua Bengio, Aaron Courville, and Yoshua Bengio. *Deep learning*, volume 1. MIT Press, 2016.
- Geoffrey E. Hinton, Simon Osindero, and Yee Whye Teh. A fast learning algorithm for deep belief nets. *Neural Computation*, 18:1527–1554, 2006.

A APPENDIX

You may include other additional sections here.